

COMP5426

Parallel and Distributed Computing

Performance Analysis of Parallel
Systems (Examples)

Jacobi Iteration


- ◆ Jacobi iteration is a numerical method used to solve Laplace partial differential equations, e.g., to determine the steady-state temperature on some domain when the temperature of its boundaries is held fixed.
- ◆ The method approaches a solution iteratively; in each iteration, the temperature of a point is computed to be the average of temperatures of its (four) neighbors, except that temperatures at the boundary are not changed.
- ◆ This iterative method is often referred to as relaxation method as an initial guess at the solution is allowed to slowly relax towards the true solution, reducing the errors as it does so.

Sequential Algorithm

- ◆ To find the solution for a two-dimensional Laplace equation:
 1. Initialise u_j to some initial *guess*.
 2. Apply the boundary conditions.
 3. For each internal mesh point set
 - ◆ $u_{ij}^* = (u_{(i+1)j} + u_{(i-1)j} + u_{i(j+1)} + u_{i(j-1)})/4$.
 4. Replace old solution U with new estimate U^* .
 5. If solution does not satisfy tolerance, repeat from step 2.





Heart of Sequential Algorithm

- ◆ Create a 2D grid and
 - ◆ Each grid point represents value of state solution at particular (i, j) location in plate
- $$u^*(i,j) = (u(i-1,j) + u(i+1,j) + u(i,j-1) + u(i,j+1)) / 4.0;$$



Parallel Algorithm 1





- ◆ Associate primitive task with each matrix element
- ◆ Agglomerate tasks in contiguous rows (rowwise block striped decomposition)



16 × 16 grid divided
among 4 processors

Communication





- ◆ Values in black cells cannot be computed without access to values held by other Associate primitive tasks.



16 × 16 grid divided
among 4 processors

Communication





- ◆ Values in black cells cannot be computed without access to values held by other Associate primitive tasks.



16 × 16 grid divided
among 4 processors

Ghost Points





- ◆ Memory locations used to store redundant copies of data held by neighboring processes
- ◆ Allocating ghost points as extra columns simplifies parallel algorithm by allowing same loop to update all cells



16 × 16 grid divided
among 4 processors

Ghost Points


- ◆ Memory locations used to store redundant copies of data held by neighboring processes
- ◆ Allocating ghost points as extra columns simplifies parallel algorithm by allowing same loop to update all cells



16 × 16 grid divided
among 4 processors

Comm in a Iteration

- ◆ At the beginning of each iteration, boundary rows are send/rcv between left and right neighbors



16 × 16 grid divided
among 4 processors

Complexity Analysis

- ◆ Sequential time complexity:
 $\Theta(n^2)$ each iteration
- ◆ Parallel computational complexity:
 $\Theta(n^2 / p)$ each iteration
- ◆ Parallel communication complexity:
 $\Theta(n)$ each iteration (two sends and two receives of n elements)


Complexity Analysis

- ◆ Sequential time: $\Theta(n^2)$ (T_s)
- ◆ Parallel time: $\Theta(n^2/p + n)$ (T_p)
- ◆ Total parallel time: $\Theta(n^2 + pn)$ (pT_p)
- ◆ Total overhead: $\Theta(pn)$ ($pT_p - T_s$)
- ◆ Speedup: $\Theta(p / (1 + p/n))$ (T_s / T_p)
- ◆ Efficiency: $\Theta(1 / (1 + p/n))$ (T_s / pT_p)
- ◆ Isoefficiency relation: ($T_s \geq CT_o$)
 $n^2 \geq Cnp \Rightarrow n \geq Cp$

Parallel Algorithm 2

- ◆ Associate primitive task with each matrix element
- ◆ Agglomerate tasks into blocks that are as square as possible (checkerboard block decomposition)
- ◆ Add rows and columns of ghost points to all four sides of rectangular region controlled by each process

Example Decomposition



16 × 16 grid divided among 16 processors

Implementation Details

- ◆ Using ghost points around 2-D blocks requires extra copying steps
- ◆ Ghost points for left and right sides are not in contiguous memory locations
 - An auxiliary buffer may be used when receiving these ghost point values and similarly, buffer may be used when sending column of values to a neighboring process

Complexity Analysis

- ◆ Sequential time: $\Theta(n^2)$
- ◆ Parallel computational time: $\Theta(n^2 / p)$
- ◆ Parallel communication time: $\Theta(n / \sqrt{p})$
each iteration (four sends and four
receives of n / \sqrt{p} elements each)

Complexity Analysis

- ◆ Sequential time: $\Theta(n^2)$
- ◆ Parallel time: $\Theta(n^2/p + n/\sqrt{p})$
- ◆ Total parallel time: $\Theta(n^2 + \sqrt{pn})$
- ◆ Total overhead: $\Theta(\sqrt{pn})$
- ◆ Speedup: $\Theta(p / (1 + \sqrt{p} / n))$
- ◆ Efficiency: $\Theta(1 / (1 + \sqrt{p} / n))$
- ◆ Isoefficiency relation: $(T_s \geq CT_o)$
 $n^2 \geq Cn\sqrt{p} \Rightarrow n \geq C\sqrt{p}$

Questions

- ◆ Complexity analysis for
 - matrix multiplication
 - odd-even transposition sort